

In [189]:

```
# House Price Prediction using Machine Learning
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import r2_score
from sklearn.neighbors import KNeighborsRegressor
from sklearn.svm import SVR
```

In [132]:

```
# Importing data
df=pd.read_csv("C:/Users/ISHITA GUPTA/Desktop//data.csv")
```

In [133]:

```
df.head()
```

Out[133]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	vie
0	2014-05-02 00:00:00	313000.00000	3.00000	1.50000	1340	7912	1.50000	0	
1	2014-05-02 00:00:00	2384000.00000	5.00000	2.50000	3650	9050	2.00000	0	
2	2014-05-02 00:00:00	342000.00000	3.00000	2.00000	1930	11947	1.00000	0	
3	2014-05-02 00:00:00	420000.00000	3.00000	2.25000	2000	8030	1.00000	0	
4	2014-05-02 00:00:00	550000.00000	4.00000	2.50000	1940	10500	1.00000	0	

In [134]:

```
df.tail()
```

Out[134]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
4595	2014-07-09 00:00:00	308166.66667	3.00000	1.75000	1510	6360	1.00000	0
4596	2014-07-09 00:00:00	534333.33333	3.00000	2.50000	1460	7573	2.00000	0
4597	2014-07-09 00:00:00	416904.16667	3.00000	2.50000	3010	7014	2.00000	0
4598	2014-07-10 00:00:00	203400.00000	4.00000	2.00000	2090	6630	1.00000	0
4599	2014-07-10 00:00:00	220600.00000	3.00000	2.50000	1490	8102	2.00000	0

In [135]:

```
df.describe()
```

Out[135]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	water
count	4600.00000	4600.00000	4600.00000	4600.00000	4600.00000	4600.00000	4600.0
mean	551962.98847	3.40087	2.16082	2139.34696	14852.51609	1.51207	0.0
std	563834.70255	0.90885	0.78378	963.20692	35884.43614	0.53829	0.0
min	0.00000	0.00000	0.00000	370.00000	638.00000	1.00000	0.0
25%	322875.00000	3.00000	1.75000	1460.00000	5000.75000	1.00000	0.0
50%	460943.46154	3.00000	2.25000	1980.00000	7683.00000	1.50000	0.0
75%	654962.50000	4.00000	2.50000	2620.00000	11001.25000	2.00000	0.0
max	26590000.00000	9.00000	8.00000	13540.00000	1074218.00000	3.50000	1.0

In [136]:

```
df.columns
```

Out[136]:

```
Index(['date', 'price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot',  
      'floors', 'waterfront', 'view', 'condition', 'sqft_above',  
      'sqft_basement', 'yr_built', 'yr_renovated', 'street', 'city',  
      'statezip', 'country'],  
      dtype='object')
```

In [137]:

```
df.shape
```

Out[137]:

```
(4600, 18)
```

In [138]:


```
pd.set_option('display.float_format', lambda x: '%.5f' % x)
```

In [139]:

```
df.describe()
```

Out[139]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	water
count	4600.00000	4600.00000	4600.00000	4600.00000	4600.00000	4600.00000	4600.0
mean	551962.98847	3.40087	2.16082	2139.34696	14852.51609	1.51207	0.0
std	563834.70255	0.90885	0.78378	963.20692	35884.43614	0.53829	0.0
min	0.00000	0.00000	0.00000	370.00000	638.00000	1.00000	0.0
25%	322875.00000	3.00000	1.75000	1460.00000	5000.75000	1.00000	0.0
50%	460943.46154	3.00000	2.25000	1980.00000	7683.00000	1.50000	0.0
75%	654962.50000	4.00000	2.50000	2620.00000	11001.25000	2.00000	0.0
max	26590000.00000	9.00000	8.00000	13540.00000	1074218.00000	3.50000	1.0



In [140]:

```
# Distinct name of cities and their count
df['city'].value_counts()
```

Out[140]:

Seattle	1573
Renton	293
Bellevue	286
Redmond	235
Issaquah	187
Kirkland	187
Kent	185
Auburn	176
Sammamish	175
Federal Way	148
Shoreline	123
Woodinville	115
Maple Valley	96
Mercer Island	86
Burien	74
Snoqualmie	71
Kenmore	66
Des Moines	58

In [141]:

```
#check null values
df.isnull().sum()
```

Out[141]:

date	0
price	0
bedrooms	0
bathrooms	0
sqft_living	0
sqft_lot	0
floors	0
waterfront	0
view	0
condition	0
sqft_above	0
sqft_basement	0
yr_built	0
yr_renovated	0
street	0
city	0
statezip	0
country	0

dtype: int64

In []:

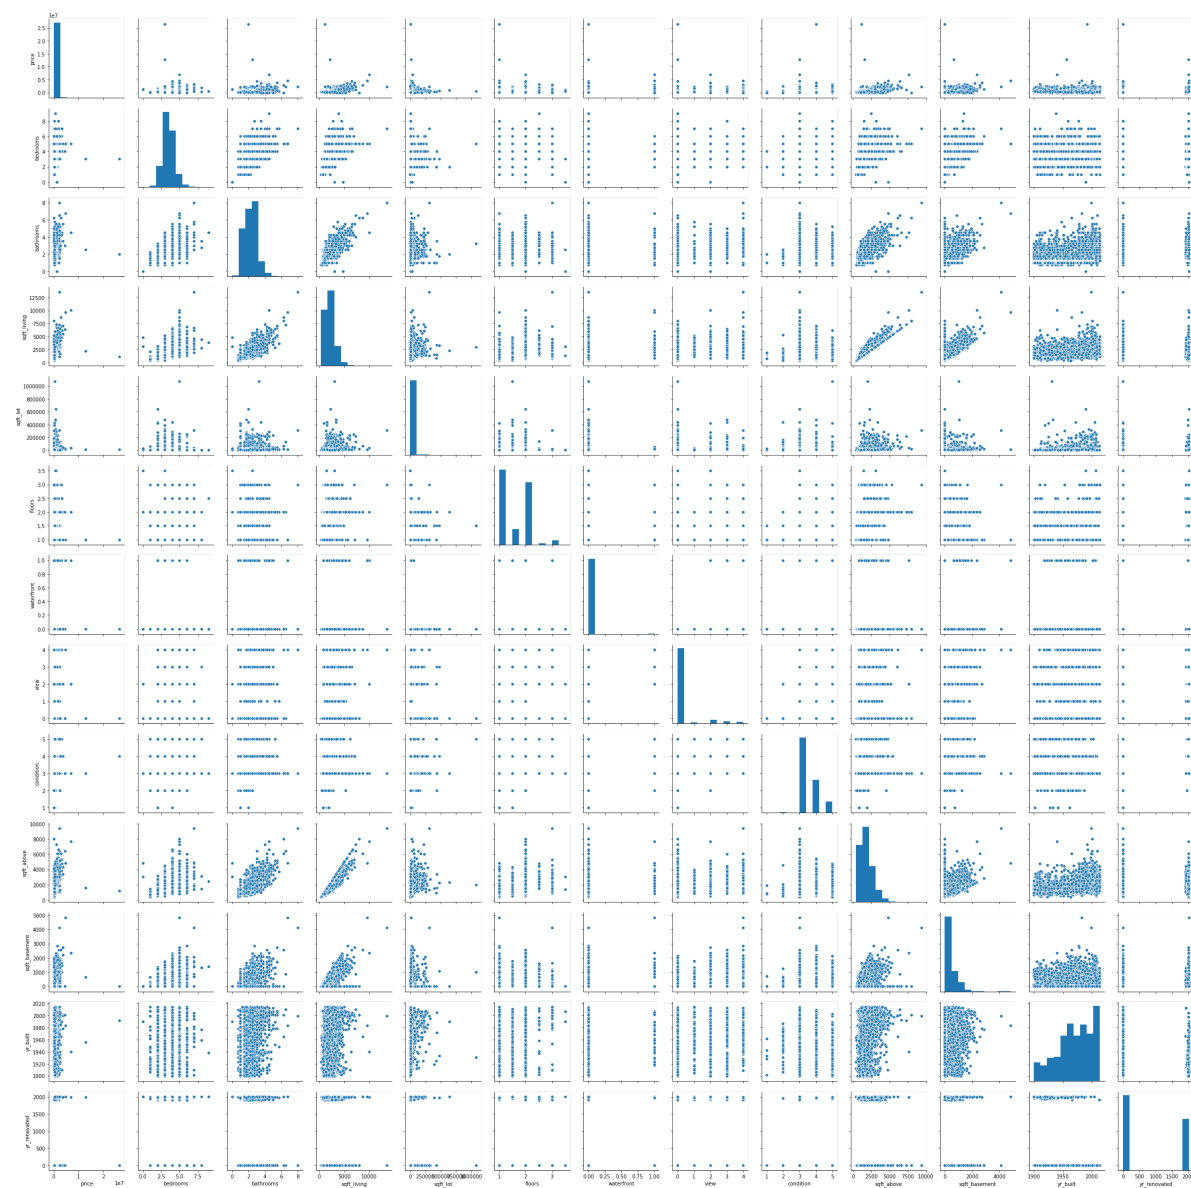
```
# Exploratory Data Analysis
```

In [142]:

```
#Plotting Data
sns.pairplot(df)
```

Out[142]:

<seaborn.axisgrid.PairGrid at 0x25373eeeeef0>




In [143]:

```
# constructing a heatmap to understand the correlation
plt.figure(figsize=(10,10))
sns.heatmap(df.corr(), cbar=True, square=True, fmt='.1f', annot=True, annot_kws={'size':8},
```

Out[143]:

<matplotlib.axes._subplots.AxesSubplot at 0x2537bddf358>



In [144]:

```
# Splitting data for training and testing
X = df[['sqft_living', 'condition', 'sqft_above', 'yr_built', 'floors']]
y = df['price']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=101)
```

In [145]:

X_train

Out[145]:

	sqft_living	condition	sqft_above	yr_built	floors
695	3200	3	3200	2004	2.00000
1170	2350	4	1430	1977	1.00000
684	700	4	700	1949	1.00000
2490	1150	3	1150	1950	1.00000
2882	1290	4	1290	1906	2.00000
979	2850	4	1450	1970	1.00000
2601	1960	4	1960	1967	1.00000
4365	1600	3	1600	2013	2.00000
3627	3500	3	3500	2005	2.00000
19	1180	3	1180	1983	1.00000

In [146]:

X_test

Out[146]:

	sqft_living	condition	sqft_above	yr_built	floors
4032	5430	4	5430	1987	2.00000
1558	2270	3	1380	1977	1.00000
2004	2070	3	2070	2004	2.00000
3186	2550	5	1860	1902	2.00000
4176	1460	3	1460	1952	1.00000
3643	930	2	930	1942	1.00000
1970	2340	5	1170	1917	1.00000
2433	1070	3	1070	1999	2.00000
1407	1930	4	1030	1967	1.00000
2042	2030	3	1720	2001	2.00000

In [147]:

```
y_train
```

Out[147]:

```
695      707000.00000
1170     555000.00000
684      267800.00000
2490     129000.00000
2882     549000.00000
979      1275000.00000
2601     285500.00000
4365     444845.00000
3627     780000.00000
19       275000.00000
1168     150000.00000
469      842500.00000
1365     292050.00000
3756     600000.00000
1909     418500.00000
3626     835000.00000
4205     410000.00000
2439     589000.00000
```

In [148]:

```
y_test
```

```
1694     809000.00000
1292     606000.00000
1439     840000.00000
2685     493000.00000
524      1300000.00000
355      530000.00000
2919     245000.00000
2466     380000.00000
699      450000.00000
178      379900.00000
346      295000.00000
2775     175000.00000
1483     380000.00000
2725     1325000.00000
1544     672000.00000
3143     715000.00000
1105     220000.00000
3797     535000.00000
166      425000.00000
154      609000.00000
```

In [149]:

```
# Standardization
std=StandardScaler()
```

In [150]:

```
# Normalization
X_train_std=std.fit_transform(X_train)
X_test_std=std.transform(X_test)
```


In [151]:

X_train

Out[151]:

	sqft_living	condition	sqft_above	yr_built	floors
695	3200	3	3200	2004	2.00000
1170	2350	4	1430	1977	1.00000
684	700	4	700	1949	1.00000
2490	1150	3	1150	1950	1.00000
2882	1290	4	1290	1906	2.00000
979	2850	4	1450	1970	1.00000
2601	1960	4	1960	1967	1.00000
4365	1600	3	1600	2013	2.00000
3627	3500	3	3500	2005	2.00000
40	1180	2	1180	1982	1.00000

In [152]:

X_train_std

Out[152]:

```
array([[ 1.0980722 , -0.66150685,  1.60694763,  1.12296229,  0.88391385],
       [ 0.21123019,  0.80496616, -0.47145962,  0.21714364, -0.96287057],
       [-1.51028665,  0.80496616, -1.32865582, -0.72222384, -0.96287057],
       ...,
       [ 0.31556455,  0.80496616,  0.72626659,  0.25069248, -0.96287057],
       [-0.17272023, -0.66150685,  0.17672163,  1.12296229,  0.88391385],
       [ 0.31556455,  0.80496616, -0.69456548, -2.36611693, -0.96287057]])
```

In [153]:

X_test_std

Out[153]:

```
array([[ 3.42472829e+00,  8.04966165e-01,  4.22550590e+00,
        5.52632029e-01,  8.83913845e-01],
       [ 1.27762709e-01, -6.61506848e-01, -5.30171685e-01,
        2.17143643e-01, -9.62870570e-01],
       [-8.09059991e-02, -6.61506848e-01,  2.80054867e-01,
        1.12296229e+00,  8.83913845e-01],
       ...,
       [-6.54744945e-01,  2.27143918e+00, -3.65777892e-01,
        -1.18344744e-01, -9.62870570e-01],
       [ 2.56148399e-03, -6.61506848e-01,  3.73994178e-01,
        4.85534352e-01,  8.83913845e-01],
       [ 8.68536621e-01, -6.61506848e-01,  1.34861452e+00,
        9.88766931e-01,  8.83913845e-01]])
```

In [154]:

```
y_train
```

```
1273    440000.00000
3912    815000.00000
1580    672500.00000
2107    378500.00000
2931    480000.00000
1949    540000.00000
4467    1337044.20000
2184    740000.00000
1530    535000.00000
49       838000.00000
4573    584000.00000
908     110700.00000
3182    625000.00000
3829    599950.00000

2623    325000.00000
973     385000.00000
4079    513000.00000
4171    749950.00000
599     450000.00000
```

In [111]:

```
y_test
```

Out[111]:

```
4032    1360000.00000
1558     332000.00000
2004     343000.00000
3186     660000.00000
4176     310000.00000
3643     100000.00000
1970     640000.00000
2433     312500.00000
1407     268000.00000
2042     471000.00000
3329     249000.00000
4589     182805.00000
3624     265000.00000
3754     660000.00000
4157     380000.00000
1124     700000.00000
2692     675000.00000
338      1039000.00000
```

In [155]:

```
# Model Training
lm = LinearRegression()
```

In [156]:

```
#Fitting data in the model  
lm.fit(X_train,y_train)
```

Out[156]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

In [157]:

```
print(lm.intercept_)
```

5335032.63650605

In [158]:


```
predictions = lm.predict(X_test)
```

In [159]:

```
# Scatter plot  
plt.scatter(y_test,predictions)
```


Out[159]:

<matplotlib.collections.PathCollection at 0x2537d6f7c88>



In [160]:

```
# Plotting Density Distribution
sns.distplot((y_test-predictions),bins=50);
```



In [161]:

```
# Calculating Errors
print('MAE:', metrics.mean_absolute_error(y_test, predictions))
print('MSE:', metrics.mean_squared_error(y_test, predictions))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

```
MAE: 189824.2263705201
MSE: 535920016816.23615
RMSE: 732065.5823191226
```

In [162]:

```
# Scoring provided by linear regression model
model_score = cross_val_score(estimator=LinearRegression(),X=X_train, y=y_train, cv=6)
model_score
```

Out[162]:

```
array([0.51925241, 0.50533965, 0.48512689, 0.55004089, 0.43191235,
       0.48796952])
```

In [167]:

```
# Using GBM algorithm(Gradient Boosting Machine)
model = GradientBoostingRegressor(n_estimators = 400, max_depth = 5, min_samples_split = 2,
                                  learning_rate = 0.1, loss = 'ls')
model.fit(X_train,y_train)

y_pred = model.predict(X_test)
model.score(X_test,y_test)
```

Out[167]:

```
0.04100223827252125
```

In [184]:

```
# Function to calculate root mean square error
def rmse (y_true,y_pred):
    return np.sqrt(np.mean((y_true - y_pred)**2))
```

In [185]:

```
# Trains model and evaluate model on test data
def model_test(model):
    model.fit(X,y)
    predictions = model.predict(X_test)
    model_rsme = rmse(y_test,predictions)
    return model_rsme
```

In [187]:

```
# K-Nearest Number
knn = KNeighborsRegressor(n_neighbors=5)
knn_rsme = model_test(knn)
print(knn_rsme)
```

609133.8329107198

In [190]:

```
# Support Vector Machine
svm = SVR()
svm_rsme = model_test(svm)
print(svm_rsme)
```

C:\Users\ISHITA GUPTA\Anaconda3\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
"avoid this warning.", FutureWarning)

774286.1715799351

In [191]:

```
# Decision Tree
tree = DecisionTreeRegressor()
tree_rsme = model_test(tree)
print(tree_rsme)
```

21795.201863224007